REINFORCEMENT LEARNING

AGENTIC AI LECTURE 8

LATEST IN AI

I WANT IT TO CONTROL MY WHOLE COMPUTER!!

manus

The general AI agent

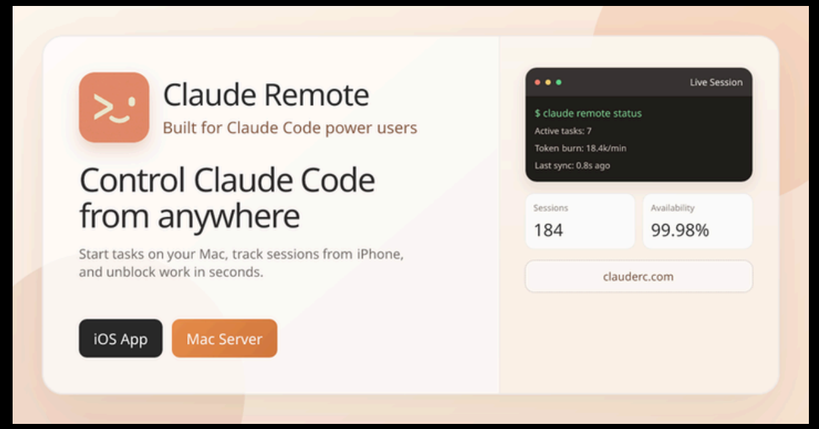

OPENCLAW

i heard about a guy in a small town in england who turned his openclaw into a short form video marketing machine

millions of views, steady app downloads, and revenue coming in every day

Bo Wang @BoWang87 · Mar 7

Few people know how popular **@openclaw** is in **China**.

Two scenes that went viral in Chinese social medias this week:

🧓 Thousands of elderly people lined up so Tencent engineers could help them install it.

🎓 A Beijing school deploying AI agents for every student.

From
Show more

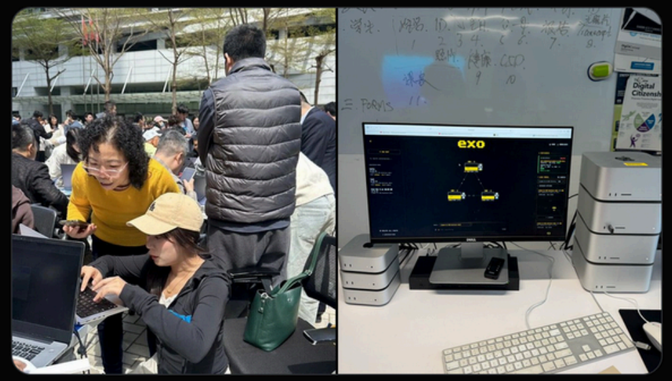

Min Choi @minchoi · Mar 24

RIP **OpenClaw** 💀

Claude now has

> Voice mode
> Agent Teams
> 38+ Connectors
> Cowork Projects
> Scheduled tasks
> Plugin Marketplace
> Persistent memory
> 1M Context window
> Dispatch for remote control
> Channels for Telegram & Discord
> Claude can use computer to run apps 💀

Agent Development Kit

Introducing GPT-5.4

I'm sharing some important news today. I have made the incredibly difficult decision to reduce the size of our team by ~10% (or ~1,600 employees). **Every**

I believe this is the right decision for Atlassian. But that doesn't mean it's easy. Far from it. I know this has a huge impact on each of you, and it weighs heavily on me and Atlassian today.

We are doing this to self-fund further investment in AI and enterprise sales, while strengthening our financial profile. We're also changing the way we work and reorganising around our System of Work to move faster.

Artificial Intelligence - AI

Securing Enterprise Agents with NVIDIA OpenShell and Cisco AI Defense

5 min read

Vikram Varakantam, Ruchika Pandey

NemoClaw

Run any agent more safely. Shape its access not its capabilities, and help keep inference private.

agents | open source | sandbox | security

What is NVIDIA NemoClaw?

reinforcment learning

what’s the point

agents make mistakes → we correct them → agents make mistakes

This process of correcting them can be automated through reinforcment learning

Trading bots
Robotic hardware movements
Sales optimization
Other autonomous agents

Here's one way to think about it - you can't beat **Djokovic** at tennis by giving an adult an entire cabinet of tennis rules, tricks of the trades and methods. You beat Djokovic by raising a child who has played tennis since he was 5 and grown up being obsessed with tennis, re-wiring his entire brain to be excellent at one thing. **THAT** is specialization. Ever realized that world champions have done what they've been doing since they were a child?

reinforcement learning

high level overview

Pre-LLM RL is different from LLM RL

Old RL

Limited environment and actions.

Adjusts model weights based on actions and resulting states, e.g. a chessboard, and moving chess pieces

New RL

LLM acts as the brain

Adjusts prompts/pipline mathematically

Tools like **https://dspy.ai/** and **https://github.com/SylphAI-Inc/AdalFlow**

in practice

1. Configure LLM
2. Prepare labeled data - create input/output examples, split into train/dev sets (like the usual ML training)
3. Define a metric/evaluation
4. Define a module - declare the input/output signature e.g. "question → answer” or “text → summary”
5. Evaluate the baseline - run the unoptimized module to get a starting score
6. Optimize - the optimizer tries different few-shot example combos and reasoning traces, scores each against your metric, keeps the best
7. Evaluate the optimized version - measure the improvement

Memento - 2000

lets look at some code

I lied

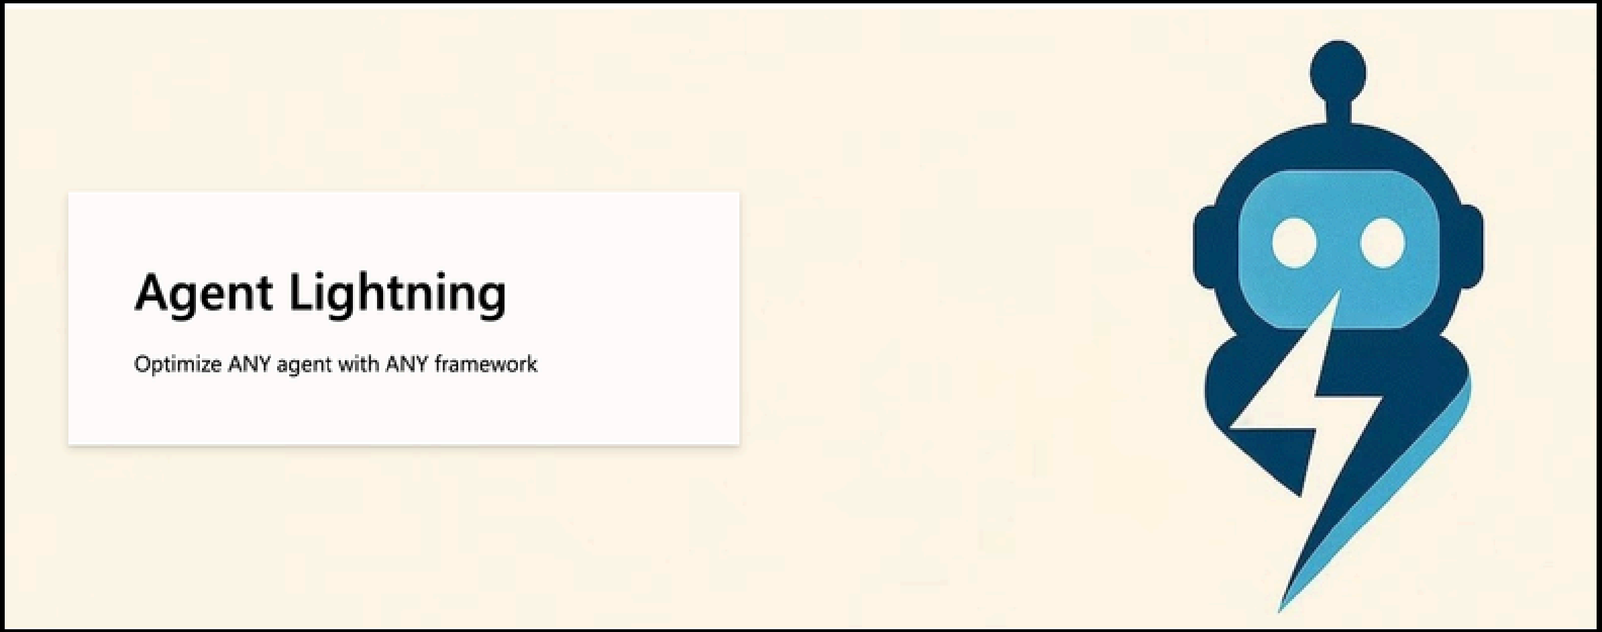

your turn